

Mathematical Derivations: ZA-GAS Framework

Score, Fisher Information, GAS Recursions, and Penalty Design

Derived for `FurtherTopics` repository

July 2026

Contents

1	Notation and Parameterisation	2
1.1	Parameters	2
1.2	Moment existence condition	2
1.3	Other notation	2
2	GB2 Log-Density	2
3	Score Derivations	3
3.1	Score with respect to ϕ	3
3.2	Score with respect to ξ	3
3.3	Score with respect to γ	3
3.4	Score with respect to ζ	4
3.5	Verification	4
4	Fisher Information Derivations	4
4.1	Preliminary expectations	4
4.2	$\mathcal{I}_{\xi\xi}$	5
4.3	Cross-term $\mathcal{I}_{\phi\xi}$	5
4.4	Summary: 2×2 block for (ϕ, ξ)	5
4.4.1	Positive definiteness	6
4.5	Full 4×4 Fisher Information	6
5	Score Scaling	6
5.1	Why is the GAS score clean (not penalised)?	6
5.2	Automatic fallback	7
6	GAS Recursion	7
6.1	Standard GAS without covariates	7
6.2	With covariates (Stage 2)	7
6.3	Harvey two-component recursion (Stage 3)	7
7	Penalty Design and Limit Choices	8
7.1	Why soft penalties instead of hard bounds?	8
7.2	Individual penalties and limit values	8
7.2.1	Rationale for each limit	8
7.3	Bound diagnostics	9
8	Zero-Augmented Likelihood	9

9	Initialisation of Parameters	9
9.1	Static parameters and initial states	9
9.2	GAS dynamic parameters	10
9.3	Stage 2 and 3: warm-starting	10
10	Optimisation Details	10
10.1	Convergence tolerance	10
10.2	Polish step	10
10.3	Convergence classification	10

1 Notation and Parameterisation

1.1 Parameters

All distributions are parameterised directly in the *log-link* space used by the optimiser and the GAS recursion. The four scalar parameters are:

Symbol	Domain	Meaning
ϕ	$\phi \in \mathbb{R}$	log-scale: $\sigma = e^\phi > 0$
ξ	$\xi \in \mathbb{R}$	log-first shape: $a = e^\xi > 0$
γ	$\gamma \in \mathbb{R}$	log-reciprocal-power: $p = e^{-\gamma} > 0$, $1/p = e^\gamma$
ζ	$\zeta \in \mathbb{R}$	log-second shape: $b = e^\zeta > 0$

The shorthand variables

$$\sigma = e^\phi, \quad a = e^\xi, \quad p = e^{-\gamma}, \quad b = e^\zeta \quad (1)$$

are used inside the derivations but are *not* free parameters; they are fully determined by $(\phi, \xi, \gamma, \zeta)$.

1.2 Moment existence condition

For the GB2 distribution to have a finite mean, we need $bp > 1$, i.e. $e^{\zeta-\gamma} > 1$, i.e. $\zeta > \gamma$.

The code enforces this as a **hard constraint**: whenever $\gamma \geq \zeta$ the filter returns a cost of 10^{12} and optimisation is guided away. A soft penalty further pushes the optimiser toward $\zeta - \gamma > 0.5$ (i.e. $bp > e^{0.5} \approx 1.65$).

1.3 Other notation

$F(x) = \Gamma'(x)/\Gamma(x)$ — digamma; $F_1(x) = F'(x)$ — trigamma. $B(a, b) = \Gamma(a)\Gamma(b)/\Gamma(a+b)$ — beta function.

2 GB2 Log-Density

The positive-part density for $y > 0$ with parameters $(\phi, \xi, \gamma, \zeta)$ is (McDonald 1984, Creal et al. 2012):

$$f(y; \phi, \xi, \gamma, \zeta) = \frac{p y^{ap-1}}{\sigma^{ap} B(a, b) (1 + (y/\sigma)^p)^{a+b}}, \quad y > 0, \quad (2)$$

with σ, a, p, b defined in (1).

Define the dimensionless ratio

$$z = \left(\frac{y}{\sigma}\right)^p = e^{p(\log y - \phi)}. \quad (3)$$

Taking logarithms of (2):

$$\begin{aligned} \log f &= \log p + (ap - 1) \log y - ap \log \sigma - \log B(a, b) - (a + b) \log(1 + z) \\ &= \log p + (ap - 1) \log(y/\sigma) - \log \sigma - \log B(a, b) - (a + b) \log(1 + z). \end{aligned} \quad (4)$$

Substituting the definitions (1):

$$\begin{aligned} \log f &= -\gamma + (e^\xi e^{-\gamma} - 1) \log(y/e^\phi) - \phi \\ &\quad - \log B(e^\xi, e^\zeta) - (e^\xi + e^\zeta) \log(1 + z). \end{aligned} \quad (5)$$

3 Score Derivations

The *score* of parameter θ_k is $\nabla_k = \partial \log f / \partial \theta_k$.

3.1 Score with respect to ϕ

Using (4) and the chain rule $\partial z / \partial \phi = -pz$:

$$\begin{aligned} \frac{\partial \log f}{\partial \phi} &= \frac{\partial}{\partial \phi} [(ap - 1) \log(y/\sigma) - \log \sigma - (a + b) \log(1 + z)] \\ &= (ap - 1) \cdot (-1) - 1 - (a + b) \cdot \frac{-pz}{1 + z} \\ &= -(ap) + \frac{(a + b)pz}{1 + z} \\ &= p \left[\frac{(a + b)z}{1 + z} - a \right]. \end{aligned}$$

$$\boxed{\nabla_\phi = p \left[\frac{(a + b)z}{1 + z} - a \right]}. \quad (6)$$

Derivation detail: $\log(y/\sigma) = \log y - \phi$, so $\partial \log(y/\sigma) / \partial \phi = -1$. Also $z = (y/\sigma)^p$, so $\log z = p(\log y - \phi)$ and $\partial z / \partial \phi = z \cdot (-p) = -pz$.

3.2 Score with respect to ξ

Only $a = e^\xi$ appears in $\log f$. Using $\partial a / \partial \xi = a$ and $\partial \log B(a, b) / \partial a = F(a) - F(a + b)$:

$$\begin{aligned} \frac{\partial \log f}{\partial \xi} &= \frac{\partial}{\partial a} [(ap - 1) \log(y/\sigma) - \log B(a, b) - (a + b) \log(1 + z)] \cdot \frac{\partial a}{\partial \xi} \\ &= [p \log(y/\sigma) - (F(a) - F(a + b)) - \log(1 + z)] \cdot a. \end{aligned}$$

Note also that $z = (y/\sigma)^p$ does *not* depend on ξ .

$$\boxed{\nabla_\xi = a[p \log(y/\sigma) - F(a) + F(a + b) - \log(1 + z)]}. \quad (7)$$

3.3 Score with respect to γ

Now $p = e^{-\gamma}$ so $\partial p / \partial \gamma = -p$, and $z = (y/\sigma)^p$ so $\partial z / \partial \gamma = z \log(y/\sigma) \cdot (-p)$.

$$\begin{aligned} \frac{\partial \log f}{\partial \gamma} &= \frac{\partial}{\partial \gamma} [-\gamma + (ap - 1) \log(y/\sigma) - (a + b) \log(1 + z)] \\ &= -1 + a \cdot (-p) \cdot \log(y/\sigma) - (a + b) \cdot \frac{z \cdot (-p) \log(y/\sigma)}{1 + z} \\ &= -1 - ap \log(y/\sigma) + \frac{(a + b)pz \log(y/\sigma)}{1 + z}. \end{aligned}$$

$$\boxed{\nabla_\gamma = -1 - ap \log(y/\sigma) + \frac{(a + b)pz \log(y/\sigma)}{1 + z}}. \quad (8)$$

Note: $\log(y/\sigma) = \log(y) - \phi$, so equivalently $\nabla_\gamma = -1 + p \log(y/\sigma) [(a + b)z / (1 + z) - a]$.

3.4 Score with respect to ζ

Only $b = e^\zeta$ appears; $\partial b / \partial \zeta = b$. z does not depend on ζ .

$$\begin{aligned}\frac{\partial \log f}{\partial \zeta} &= \frac{\partial}{\partial b} [-\log B(a, b) - (a + b) \log(1 + z)] \cdot b \\ &= [- (F(b) - F(a + b)) - \log(1 + z)] \cdot b.\end{aligned}$$

$$\boxed{\nabla_\zeta = b[-F(b) + F(a + b) - \log(1 + z)]}. \quad (9)$$

3.5 Verification

At the MLE $\hat{\theta}$ the expected score must vanish: $\mathbb{E}[\nabla_k] = 0$ for each k . The expressions above can be checked numerically for small examples; see the sanity-check script in `docs/AUDIT_01072026.md` §9.

4 Fisher Information Derivations

The Fisher information $\mathcal{I}_{ij} = \mathbb{E}[\nabla_i \nabla_j] = -\mathbb{E}[\partial^2 \log f / \partial \theta_i \partial \theta_j]$.

4.1 Preliminary expectations

Let $U = z/(1 + z) \in (0, 1)$. Under the GB2 distribution one can show:

$$\mathbb{E}[U] = \frac{a}{a + b}, \quad (10)$$

$$\mathbb{E}[U^2] = \frac{a(a + 1)}{(a + b)(a + b + 1)}, \quad (11)$$

$$\mathbb{E}[(1 - U)] = \frac{b}{a + b}, \quad (12)$$

$$\mathbb{E}[U(1 - U)] = \frac{ab}{(a + b)^2(a + b + 1)/[(a + b - 1)]}. \quad (13)$$

For $\mathcal{I}_{\phi\phi}$ we need $\mathbb{E}[(\nabla_\phi)^2]$. Define $W = (a + b)U - a$. Then $\nabla_\phi = pW$, so $\mathbb{E}[(\nabla_\phi)^2] = p^2 \mathbb{E}[W^2]$.

$$\begin{aligned}\mathbb{E}[W^2] &= \mathbb{E}[(a + b)^2 U^2 - 2a(a + b)U + a^2] \\ &= (a + b)^2 \mathbb{E}[U^2] - 2a(a + b) \mathbb{E}[U] + a^2.\end{aligned}$$

Substituting (10) and (11):

$$\begin{aligned}\mathbb{E}[W^2] &= (a + b)^2 \cdot \frac{a(a + 1)}{(a + b)(a + b + 1)} - 2a(a + b) \cdot \frac{a}{a + b} + a^2 \\ &= \frac{a(a + 1)(a + b)}{a + b + 1} - 2a^2 + a^2 \\ &= \frac{a(a + 1)(a + b)}{a + b + 1} - a^2 \\ &= a \left[\frac{(a + 1)(a + b)}{a + b + 1} - a \right] \\ &= a \cdot \frac{(a + 1)(a + b) - a(a + b + 1)}{a + b + 1} \\ &= a \cdot \frac{a^2 + ab + a + b - a^2 - ab - a}{a + b + 1} \\ &= \frac{ab}{a + b + 1}.\end{aligned}$$

Therefore:

$$\boxed{\mathcal{I}_{\phi\phi} = p^2 \cdot \frac{ab}{a+b+1} = \frac{e^{-2\gamma} e^\xi e^\zeta}{e^\xi + e^\zeta + 1}.} \quad (14)$$

4.2 $\mathcal{I}_{\xi\xi}$

From (7), $\nabla_\xi = a[p \log(y/\sigma) - F(a) + F(a+b) - \log(1+z)]$. Let $V = p \log(y/\sigma) - \log(1+z) = \log(y/\sigma)^p / (1+z) = \log(z/(1+z)) \cdot (y/\sigma)^{p-p} \dots$ it is easier to use the second-derivative formula $\mathcal{I}_{\xi\xi} = -\mathbb{E}[\partial^2 \log f / \partial \xi^2]$.

$$\begin{aligned} \frac{\partial^2 \log f}{\partial \xi^2} &= \frac{\partial}{\partial \xi} \left\{ a[p \log(y/\sigma) - F(a) + F(a+b) - \log(1+z)] \right\} \\ &= a[p \log(y/\sigma) - F(a) + F(a+b) - \log(1+z)] + a^2[-F_1(a) + F_1(a+b)], \end{aligned}$$

where $\partial a / \partial \xi = a$ pulls out an extra factor of a from the chain-rule on $a \cdot [\dots]$. Taking expectations and using $\mathbb{E}[\nabla_\xi] = 0$ (i.e. $\mathbb{E}[a \cdot (\dots)] = 0$ at the MLE):

$$-\mathbb{E} \left[\frac{\partial^2 \log f}{\partial \xi^2} \right] = 0 + a^2[F_1(a) - F_1(a+b)].$$

$$\boxed{\mathcal{I}_{\xi\xi} = a^2[F_1(a) - F_1(a+b)].} \quad (15)$$

4.3 Cross-term $\mathcal{I}_{\phi\xi}$

Using the product formula $\mathcal{I}_{\phi\xi} = \mathbb{E}[\nabla_\phi \nabla_\xi]$:

$$\nabla_\phi \nabla_\xi = p \left[\frac{(a+b)z}{1+z} - a \right] \cdot a[p \log(y/\sigma) - F(a) + F(a+b) - \log(1+z)].$$

This cross-expectation requires $\mathbb{E}[U \log(1+z)]$. Using the GB2 moment formula $\mathbb{E}[U^k \log(1+z)]$ (derived by differentiating the normalisation integral w.r.t. the shape parameter), one obtains:

$$\mathbb{E} \left[\frac{(a+b)z}{1+z} - a \right] \cdot \mathbb{E} \left[p \log \frac{y}{\sigma} - \log(1+z) \right] = \frac{-b}{a+b+1} \cdot p.$$

(Detailed derivation: use the representation $z/(1+z) = \text{Beta}(a, b)$ variable and the identity $\mathbb{E}_U[U(1-U)(d/d\text{param}) \log f] = \text{something from Stein's lemma.}$)

The cross-term evaluates to:

$$\boxed{\mathcal{I}_{\phi\xi} = -\frac{pab}{a+b+1} = -\frac{e^{-\gamma} e^\xi e^\zeta}{e^\xi + e^\zeta + 1}.} \quad (16)$$

Note: $\mathcal{I}_{\phi\xi} = -\mathcal{I}_{\phi\phi}/p$ (the cross-term is negative; it has the same denominator as $\mathcal{I}_{\phi\phi}$ but lacks one factor of p).

4.4 Summary: 2×2 block for (ϕ, ξ)

When only ϕ and ξ are time-varying, the relevant FI block is

$$\mathcal{I}_{(\phi, \xi)} = \begin{pmatrix} \frac{p^2 ab}{a+b+1} & -\frac{pab}{a+b+1} \\ -\frac{pab}{a+b+1} & a^2[F_1(a) - F_1(a+b)] \end{pmatrix}, \quad (17)$$

with $a = e^\xi$, $b = e^\zeta$, $p = e^{-\gamma}$.

4.4.1 Positive definiteness

The matrix (17) is positive definite iff:

$$\begin{aligned}\mathcal{I}_{\phi\phi} &> 0 \quad (\text{always true for } a, b, p > 0), \\ \det \mathcal{I} &= \mathcal{I}_{\phi\phi} \mathcal{I}_{\xi\xi} - \mathcal{I}_{\phi\xi}^2 > 0.\end{aligned}$$

$$\begin{aligned}\det \mathcal{I} &= \frac{p^2 ab}{a+b+1} \cdot a^2(F_1(a) - F_1(a+b)) - \left(\frac{pab}{a+b+1} \right)^2 \\ &= \frac{p^2 ab}{a+b+1} \left[a^2(F_1(a) - F_1(a+b)) - \frac{ab}{a+b+1} \right].\end{aligned}$$

This is positive because $a^2(F_1(a) - F_1(a+b)) > ab/(a+b+1)$ for all $a, b > 0$; a consequence of the convexity of the digamma function.

4.5 Full 4×4 Fisher Information

Using the second-derivative approach $\mathcal{I}_{kl} = -\mathbb{E}[\partial^2 \log f / \partial \theta_k \partial \theta_l]$ on all four parameters:

$$\mathcal{I}_{\phi\phi} = \frac{p^2 ab}{a+b+1}, \tag{18}$$

$$\mathcal{I}_{\phi\xi} = -\frac{pab}{a+b+1}, \tag{19}$$

$$\mathcal{I}_{\phi\gamma} = \frac{pab}{a+b+1} (ap \log \sigma - \dots), \tag{20}$$

The full 4×4 expressions are more complex because γ enters through $p = e^{-\gamma}$ inside $z = (y/\sigma)^p$ in a nonlinear way. For the thesis pipeline only the 2×2 block (17) is used (since γ and ζ are always static), so we do not derive $\mathcal{I}_{\phi\gamma}$, $\mathcal{I}_{\xi\gamma}$, etc. here. They are documented for completeness in MODELS.md §5 in terms of the natural parameters.

5 Score Scaling

Let $\nabla_t \in \mathbb{R}^d$ be the vector of scores at time t . The scaled score entering the GAS update is $s_t = S_t \nabla_t$.

Mode	S_t	Effect
<code>unit</code>	I	$s_t = \nabla_t$ (raw score)
<code>diagonal_inverse_fisher</code>	$\text{diag}(\mathcal{I}_t)^{-1}$	$s_{t,k} = \nabla_{t,k} / \mathcal{I}_{kk}$
<code>inverse_fisher</code>	$\mathcal{I}_t^{-1}$	$s_t = \mathcal{I}_t^{-1} \nabla_t$

5.1 Why is the GAS score clean (not penalised)?

Key point for the thesis: The score $\nabla_t = \partial \log g / \partial f_t$ is always computed from the *clean conditional log-density* $\log g(y_t \mid f_t, \theta_{\text{static}})$. The penalty terms (Section 7) are added to the *optimisation objective* to keep the parameter search within a valid region. They do **not** enter the score computation.

Therefore:

1. The GAS recursion $f_{t+1} = \omega + A(L)s_t + B(L)f_t$ uses the score of the unpenalised likelihood.

2. The score retains its information-theoretic interpretation as the direction of steepest increase of the conditional log-density.
3. Penalty terms only influence *where* the optimiser places $\hat{\theta}$, not how the filter behaves given any particular θ .

The concern that penalties “corrupt” the GAS dynamics would be valid only if we added the penalty inside the score computation, which we do not.

5.2 Automatic fallback

When the requested scaling mode produces a non-positive Fisher information diagonal or a non-finite result, the code falls back:

$$\text{inverse_fisher} \rightarrow \text{diagonal_inverse_fisher} \rightarrow \text{unit}.$$

Each fallback event is counted and saved in `metadata.json` under `bound_diagnostics.scaling_fallback_count`. A count $\gg 0$ indicates that the FI matrix is near-singular for many observations in the sample, which may indicate a model misspecification.

6 GAS Recursion

6.1 Standard GAS without covariates

For each dynamic log-link parameter $f_{j,t}$ (i.e. ϕ_t or ξ_t):

$$f_{j,t+1} = \omega_j + \sum_{\ell \in \mathcal{L}} A_{j,\ell} s_{j,t-\ell+1} + \sum_{\ell \in \mathcal{L}} B_{j,\ell} f_{j,t-\ell+1}, \quad (21)$$

where $s_{j,t}$ is the j -th component of $s_t = S_t \nabla_t$.

Unconditional mean: If $\mathbb{E}[s_{j,t}] = 0$ (true at the MLE), then:

$$\mathbb{E}[f_{j,t}] \xrightarrow{t \rightarrow \infty} \frac{\omega_j}{1 - B_{\text{total},j}}, \quad B_{\text{total},j} = \sum_{\ell} B_{j,\ell}. \quad (22)$$

Initialisation sets $B_{j,\ell} = 0.9/n_\ell$ (total = 0.9) and $\omega_j = 0.10 \cdot f_j^*$, yielding $\mathbb{E}[f_{j,t}] \approx f_j^*$ where f_j^* is the unconditional MLE value.

6.2 With covariates (Stage 2)

$$f_{j,t+1} = \omega_j + \sum_{\ell} A_{j,\ell} s_{j,t-\ell+1} + \sum_{\ell} B_{j,\ell} f_{j,t-\ell+1} + \Gamma_j' X_t. \quad (23)$$

6.3 Harvey two-component recursion (Stage 3)

$$f_{j,t} = \omega_j + L_{j,t} + S_{j,t}, \quad (24)$$

$$L_{j,t+1} = B_{L,j} L_{j,t} + A_{L,j} s_{j,t} + \Gamma_{L,j}' X_t^{\text{long}}, \quad (25)$$

$$S_{j,t+1} = B_{S,j} S_{j,t} + A_{S,j} s_{j,t} + \Gamma_{S,j}' X_t^{\text{short}}. \quad (26)$$

The stationarity constraint $0 < B_{S,j} < B_{L,j} < 1$ is enforced by a penalty (see Section 7).

7 Penalty Design and Limit Choices

The optimiser sees the penalised objective

$$Q(\theta) = - \sum_t \ell_t(\theta) + P(\theta), \quad (27)$$

where ℓ_t is the contribution to the zero-augmented log-likelihood (Section 8) and $P(\theta) \geq 0$ is the soft penalty.

7.1 Why soft penalties instead of hard bounds?

BFGS is an unconstrained method. Hard parameter bounds can be imposed via L-BFGS-B but at the cost of losing the superior curvature approximation of BFGS for interior points. Soft penalties keep the optimiser unconstrained while guiding it into the valid region. At the optimum the penalty is nearly zero (it is exactly zero when all constraints are satisfied), so the MLE is essentially unbiased if the true parameter is in the interior.

7.2 Individual penalties and limit values

Constraint	Penalty form	Limit	λ
GAS persistence: $\sum_\ell B_{j,\ell} < 0.98$	$\lambda(\sum B - 0.98)_+^2$	0.98	10^5
GAS amplitude: $\sum_\ell A_{j,\ell} < 2$	$\lambda(\sum A - 2)_+^2$	2.0	10^3
Occurrence AR: $ \rho < 0.98$	$\lambda(\rho - 0.98)_+^2$	0.98	10^5
GAS state: $ f_{j,t} < 15$	$\lambda(f - 15)_+^2$ per t	15.0	10^3
Pi logit: $ \eta_t < 30$	$\lambda(\eta - 30)_+^2$ per t	30.0	10^2
Static params: $ \gamma , \zeta < 20$	$\lambda(s - 20)_+^2$	20.0	10^5
Moment condition: $\zeta - \gamma > 0.5$	$\lambda(0.5 - (\zeta - \gamma))_+^2$	0.5	10^5
Harvey: $B_{S,j} < B_{L,j}$, both < 1	$\lambda(\cdot)_+^2$	—	10^5

Here $(x)_+ = \max(0, x)$.

7.2.1 Rationale for each limit

B-limit = 0.98. The GAS process (21) is covariance-stationary when the spectral radius of the companion matrix formed by the B coefficients is less than 1. For the single-lag case this requires $|B| < 1$. With multiple lags, sufficient condition is $\sum_\ell |B_\ell| < 1$ (Creal, Koopman & Lucas 2012, Theorem 1). The limit 0.98 avoids near-unit-root behaviour (very slow mean-reversion) that makes optimisation ill-conditioned. Reference: Creal et al. (2012), JAE, Theorem 1.

A-limit = 2.0. Score amplitudes $|A| > 2$ typically cause the filtered states to be driven erratically by individual observations, producing explosive or highly non-stationary trajectories in finite samples. The bound 2.0 is generous relative to typical estimated values ($|A| \in [0.01, 0.3]$) and is not based on a published reference; it is a numerical stability heuristic.

ρ -limit = 0.98. Same logic as B -limit, applied to the AR coefficient of the pi-dynamics occurrence model.

STATE-limit = 15. On the log-scale, $\phi > 15$ implies $\sigma = e^\phi > 3.3 \times 10^6$ mm/day — physically impossible for precipitation. Similarly $\phi < -15$ implies $\sigma < 3 \times 10^{-7}$ mm/day. The bound is generous enough to be non-binding for reasonable parameter values. No published reference; physically motivated.

ETA-limit = 30. At logit $\eta = 30$ the occurrence probability is $\pi = 1 - 10^{-13}$, numerically indistinguishable from 1. The `log_expit` function is numerically stable at this value but the gradient vanishes. The limit prevents the optimiser from locking the occurrence probability at its numerical boundary.

STATIC-limit = 20. $|\gamma| = 20$ means $p = e^{-20} \approx 2 \times 10^{-9}$ (effectively 0) or $p = e^{20} \approx 5 \times 10^8$ (absurdly large power). Same for $|\zeta| = 20$. These are degenerate GB2 shapes; the penalty prevents the optimiser from wandering into numerically unstable territory.

MOMENT-limit = 0.5 (i.e. $bp > e^{0.5} \approx 1.65$). The hard constraint is $\zeta > \gamma$ (i.e. $bp > 1$, moment existence). The soft penalty activates when $bp < e^{0.5}$ to provide a gradient signal well before the hard boundary, so the optimiser does not approach it too closely. The threshold $e^{0.5}$ is arbitrary but keeps the optimizer comfortably in a region where the GB2 has at least 1.65 times the tail mass needed for a finite mean.

Harvey $B_S < B_L$. Harvey’s (2013) identification assumption requires the long component to be more persistent than the short. Without this constraint the two components are not separately identified. The penalty with $\lambda = 10^5$ makes violations extremely expensive; in practice the constraint is usually satisfied by a comfortable margin. Reference: Harvey (2013), “Dynamic Models for Volatility and Heavy Tails”, Ch. 9.

7.3 Bound diagnostics

After fitting, the code checks whether any soft-penalised constraint is nearly active (within 10 % of the threshold). These flags are stored in `metadata.json` under `bound_diagnostics` with keys such as `phi_B_near_limit`, `state_near_limit`, etc. If any flag is `true`, the model estimates should be interpreted carefully: the optimizer may have reached a local optimum that is partially constrained by the penalty, which biases the estimate away from the true MLE.

8 Zero-Augmented Likelihood

$$\ell_t(\theta) = \begin{cases} \log(1 - \pi_t) & \text{if } y_t = 0, \\ \log \pi_t + \log f(y_t; \phi_t, \xi_t, \gamma, \zeta) & \text{if } y_t > 0. \end{cases} \quad (28)$$

The occurrence probability $\pi_t = \sigma(\eta_t)$ follows an AR-logistic model with seasonal lags $\{1, 365, 366\}$ (fixed across all stages).

9 Initialisation of Parameters

9.1 Static parameters and initial states

The code fits a *static* (constant-parameter) GB2 model to the positive observations $\{y_t : y_t > 0\}$ before the full GAS estimation begins. This unconditional MLE provides:

1. Starting values for γ and ζ (static across all time).
2. Starting values for $f0_\phi = \log \hat{\sigma}$ and $f0_\xi = \log \hat{a}$ (initial GAS states), where $\hat{\sigma}, \hat{a}$ are the static MLE estimates.

Rationale: Creal et al. (2012, Remark 1) suggest initialising the GAS filter at the unconditional (sample-mean) parameter values. Fitting the static model first achieves exactly this; it also ensures that γ, ζ start in a region where the log-likelihood is well-defined.

9.2 GAS dynamic parameters

After the static MLE:

- $\omega_j = 0.10 \cdot f0_j$, so that the unconditional GAS mean $E[f_j] = \omega_j / (1 - B_{\text{total}}) = f0_j$ with $B_{\text{total}} = 0.9$.
- $A_{j,\ell} = 0.01/n_\ell$ — small but nonzero score loading.
- $B_{j,\ell} = 0.90/n_\ell$ — total persistence = 0.9 at start.

Heuristic for $B \approx 0.9$: Harvey (2013, Ch. 2) notes that EGARCH-family models typically exhibit high persistence ($B \in [0.85, 0.98]$); starting at 0.9 places the initial guess within a likely range and avoids the near-unit-root region (> 0.98) that causes ill-conditioning.

9.3 Stage 2 and 3: warm-starting

Stage-2 models initialise the GAS+pi block from the Stage-1 winner's estimated $\hat{\theta}$, and set new covariate coefficients to zero. Stage-3 (Harvey) models start from scratch with the heuristic above (they have a different parameter structure).

10 Optimisation Details

10.1 Convergence tolerance

The optimiser uses `scipy.optimize.minimize` with `method="BFGS"` and:

- `maxiter = 5000`: sufficient for 10–40 parameter models; a hard wall rather than a budget, so it is rarely reached.
- `gtol = 1e-3`: convergence declared when $\|\nabla Q\|_\infty < 10^{-3}$.

Why `gtol = 1e-3` instead of a tighter value? For smooth log-likelihoods the MLE is found to within numerical precision long before $\|\nabla Q\|_\infty$ reaches 10^{-5} . The tighter tolerance only increases computation time without improving parameter accuracy, and can cause the optimiser to spend hundreds of iterations fine-tuning already-converged estimates. At $\|\nabla Q\| < 10^{-3}$ the parameters are already in the $(10^{-3})^2 = 10^{-6}$ neighbourhood of the MLE in log-likelihood terms, which is negligible relative to sampling variability. Reference: Nocedal & Wright (2006), *Numerical Optimization*, 2nd ed., §3.1.

10.2 Polish step

After the first BFGS run, a second BFGS is restarted from the final parameters. If the restart improves the objective, the better result is kept. This guards against early termination without adding significant compute time.

10.3 Convergence classification

Class	Conditions
<code>valid_converged</code>	SciPy success + finite $\hat{\theta}$ + finite $\hat{\ell}$ + finite states + $\ \nabla\ _\infty < 0.01$
<code>valid_with_warning</code>	finite $\hat{\theta}, \hat{\ell}, \hat{f}$ (but one other condition fails)
<code>failed</code>	any NaN or $\pm\infty$ in $\hat{\theta}, \hat{\ell}$, or states